

## **MODULE - 3**

### **1. Explain the Map-reduce in NoSQL.**

**Map:** In the map phase, the input data is divided into smaller chunks and processed independently by multiple nodes in a distributed system. Each node applies a map function to the data it receives and generates intermediate key-value pairs as output.

**Shuffle:** After the map phase, the intermediate key-value pairs are shuffled and redistributed across the nodes based on the keys. This step ensures that all key-value pairs with the same key are grouped together on the same node.

**Reduce:** In the reduce phase, each node processes the shuffled intermediate data and applies a reduce function to combine the values associated with each key. The reduce function produces a final set of key-value pairs, which is the output of the Map-Reduce operation.

**Parallel Processing:** Map-Reduce leverages the power of parallel processing by allowing multiple nodes to perform computations simultaneously on different portions of the data. This parallelization enables faster processing of large datasets compared to traditional sequential approaches.

**Scalability:** Map-Reduce is designed to handle large-scale data processing efficiently. By distributing the workload across multiple nodes, it can handle massive amounts of data in a scalable manner, making it suitable for Big Data applications.

### **2. Explain the Composing Map-Reduce Calculation.**

**Break it down:** To solve a big problem, we break it down into smaller parts. Each part is called a Map, where we apply a specific action to the data.

**Grouping:** After mapping, we group together similar results based on their common properties.

**Combine and simplify:** Once the grouping is done, we combine the results by taking similar things and summarizing them. This makes the data simpler and easier to work with.

**Repeat if needed:** Sometimes, we need to do more actions on the combined data. So we go back to step 1 and apply another set of actions on the simplified results. We can repeat this step as many times as necessary.

**Final result:** After all the necessary actions are done, we take the simplified data and apply one last step of combining and summarizing. This gives us the final result of our composed Map-Reduce calculation.

In simple terms, composing Map-Reduce calculations means breaking down a big problem into smaller actions, grouping similar results, combining and simplifying them, and repeating these steps if needed. Finally, we get the desired result by applying the last set of actions on the simplified data.

### **3. Explain the key-store features.**

**Key-Value Storage:** The fundamental feature of a key-value store is the ability to store data as key-value pairs. Each piece of data (value) is associated with a unique identifier (key). This structure allows efficient and fast access to data based on its key.

**Simplicity:** Key-value stores have a simple and straightforward data model. They do not enforce any schema or complex relationships between data. This simplicity makes key-value stores easy to use and understand.

**High Performance:** Key-value stores are optimized for high performance. They are designed to handle a massive amount of data and provide fast access times.

**Scalability:** Key-value stores are built to scale horizontally, meaning they can handle increasing amounts of data by adding more servers or nodes. This allows them to handle large datasets and handle high traffic loads.

**Flexibility:** Key-value stores can handle various types of data, such as strings, numbers, and binary data. They provide flexibility in storing and retrieving different types of values associated with their corresponding keys.

#### 4. Explain the Transactions in NoSQL.

**All or Nothing:** Transactions work like a "do it all or do nothing" rule. When you perform a transaction, either all the changes you make to the data are successfully saved, or none of them are. This helps keep the data consistent and avoids leaving it in an incomplete or incorrect state.

**Keeping Things in Order:** Transactions ensure that the data follows certain rules or constraints. These rules help maintain the data's correctness and prevent any invalid or conflicting changes. If any operation breaks these rules, the whole transaction is canceled, and the data is reverted back to its original state.

**No Confusion Allowed:** Transactions make sure that different operations happening at the same time don't interfere with each other. Each transaction is treated as if it's the only one running, so there are no conflicts or confusions. This ensures that the data remains reliable and doesn't get messed up due to simultaneous changes.

**Keeping it Safe:** Once a transaction is completed, its changes are permanently saved and won't be lost even if there is a system crash or failure. This durability ensures that the data stays secure and recoverable, giving you peace of mind.

**Dependable and Reliable:** Transactions in NoSQL databases follow the ACID properties, which stand for Atomicity, Consistency, Isolation, and Durability. These properties make transactions dependable and reliable, ensuring that your data stays accurate, protected, and available when you need it.

#### 5. Explain the query features.

Query features in databases refer to the capabilities and functionalities that enable users to retrieve specific information from the stored data.

**Data Retrieval:** Query features allow users to retrieve data based on specific criteria or conditions. Users can specify what data they want to retrieve, such as specific fields, records that meet certain conditions, or data that matches specific patterns.

**Filtering and Sorting:** Query features enable users to filter and sort the retrieved data. They can apply conditions or rules to narrow down the results and only get the data that meets certain criteria. Users can also specify how the data should be sorted, such as ascending or descending order based on a specific field.

**Aggregation and Summarization:** Query features support aggregation and summarization operations on the data. Users can perform functions like calculating sums, averages, counts, or grouping data based on specific fields to obtain meaningful insights or summaries of the data.

**Joins and Relationships:** Query features allow users to combine data from multiple tables or collections by using joins or establishing relationships. This enables users to retrieve data that is related or connected across different datasets, providing a more comprehensive view of the information.

**Indexing and Performance Optimization:** Query features often include indexing mechanisms to improve the performance of data retrieval. Indexes create optimized data structures that enable faster searching and retrieval of data based on specific fields, enhancing the overall query performance.

## 6. Explain Structure data.

Structured data refers to data that is organized and formatted in a consistent and predefined manner. It follows a specific structure or schema, making it easy to understand and process.

**Organized Format:** Structured data is arranged in a well-defined format with a clear structure. It is typically represented in tabular form, like rows and columns, similar to a spreadsheet. Each piece of data has a specific place or field where it belongs.

**Schema or Blueprint:** Structured data follows a schema, which is a blueprint or set of rules that define the structure of the data. The schema specifies the type of data that can be stored, the format it should be in, and any constraints or rules that need to be followed.

**Clear Relationships:** Structured data often includes relationships between different elements. For example, in a relational database, data in one table can be related to data in another table through keys or identifiers. These relationships help establish connections and provide a comprehensive view of the data.

**Easy to Analyze and Process:** Due to its organized format and predefined structure, structured data is easier to analyze and process. It allows for efficient querying, filtering, sorting, and aggregating operations. Analysis tools and algorithms can be applied directly to structured data, making it straightforward to derive insights and perform calculations.

**Compatibility with Applications:** Structured data is highly compatible with various applications and tools. It can be seamlessly integrated with relational databases, spreadsheet software, and other data processing systems. Its structured nature allows for easy integration and interoperability with different technologies.

## 7. Explain some use case of NoSQL

**Handling Big Data:** NoSQL databases are great for dealing with really large amounts of data. They can handle and process huge volumes of information quickly and efficiently.

**Powering Websites and Blogs:** NoSQL databases work well for websites and blogs that need to manage lots of content. They can store different types of content like text, images, and videos in a flexible way. NoSQL databases are good at quickly finding and updating content, which helps in serving dynamic information on websites and blogs.

**Personalized Recommendations:** NoSQL databases are handy for creating personalized experiences for users. They can store information about people's preferences, behavior, and interests. By using this data, NoSQL databases can provide personalized recommendations for products, services, or content, making users feel like the application understands their needs.

**Internet of Things (IoT) Devices:** NoSQL databases are a good fit for managing data from IoT devices like smart sensors and gadgets. These devices generate a lot of data quickly. NoSQL databases can handle this high data flow and can scale up to support large IoT networks. They're often used to collect and process data from smart homes, factories, or connected cars.

**Managing User Sessions:** NoSQL databases are useful for keeping track of user sessions in web applications. They can store information about user preferences, shopping carts, or logged-in status.

## 8. Explain store session information

Storing session information refers to the process of keeping track of user-specific data during their interaction with a website or application.

**User Identification:** When a user visits a website or application, a unique identifier, often called a session ID, is assigned to them. This identifier helps identify and differentiate one user from another.

**Data Storage:** Session information includes data specific to each user's session, such as preferences, settings, items in a shopping cart, or temporary data related to their current activity. This information needs to be stored and associated with the user's session ID for easy retrieval.

**Session Data Management:** The session information is typically stored in a database or memory cache specifically designed for handling session data. NoSQL databases are commonly used for this purpose due to their ability to handle high-volume, fast read/write operations.

**Retrieval and Updates:** Whenever the user interacts with the website or application, their session data can be quickly retrieved based on their session ID. This allows the application to personalize the user experience and provide customized content. The session data can also be updated in real-time as the user's session progresses.

**Session Expiration:** Sessions usually have an expiration time or are automatically terminated after a period of inactivity. This helps manage server resources and maintain security. When a session expires, the associated session data is either deleted or marked as inactive.

## 9. Explain Preference in NoSQL

Preferences in NoSQL refer to the ability to define and customize how data is stored and accessed in a NoSQL database.

**Data Modeling Flexibility:** NoSQL databases offer flexibility in defining how data is structured and organized. You have the freedom to choose the data model that best suits your application's needs, such as key-value, document, columnar, or graph. This flexibility allows you to model your data based on how you want to query and retrieve it.

**Schema-less Nature:** NoSQL databases are schema-less, meaning you don't have to define a fixed schema upfront. You can store different types of data and easily add or modify fields as needed without affecting other documents or records. This flexibility allows for agile development and accommodates changing data requirements.

**Horizontal Scalability:** NoSQL databases are designed to scale horizontally by distributing data across multiple servers or nodes. This scalability enables you to handle growing amounts of data and high traffic loads. You can add more servers to the cluster to increase storage capacity and performance.

**Performance Optimization:** NoSQL databases provide various mechanisms to optimize performance. For example, you can define indexes on specific fields to

speed up data retrieval. Additionally, you can choose different consistency models (like eventual consistency) to balance performance and data freshness based on your application's requirements.

**Customized Access Patterns:** NoSQL databases allow you to design access patterns tailored to your application's specific needs.

## 10. Explain relationships among data multioperationtransactions

Relationships among data in multi-operation transactions refer to the connections and dependencies between different pieces of data that are involved in a transaction.

**Keeping Things Together:** In a multi-operation transaction, all the actions you take on the data are treated as a single group. It's like bundling them up together. This ensures that either all the actions succeed and are completed, or none of them happen at all.

**Playing by the Rules:** Transactions make sure that the relationships between different pieces of data are maintained correctly. There are rules or constraints in place that define how the data should relate to each other.

**Protecting Data's Good Name:** Relationships among data are important for keeping the data in good shape. Transactions make sure that when you make changes to the data, it doesn't end up in a messy or incorrect state.

**No Conflicts Allowed:** Transactions make sure that multiple actions happening at the same time don't clash or create problems.

**Changes that Stick:** Once a multi-operation transaction is done and confirmed, the changes made to the data become permanent and won't be lost even if something goes wrong. The changes are saved securely, so the relationships among the data remain intact even in the face of system failures or crashes.